


Machine Learning

Zusammenfassung / Joel von Rotz /  [Quelldateien](#)

Begriffe

Artificial Intelligence Software, die intelligentes Verhalten imitiert

→ **Machine Learning** System, das von Erfahrung/Daten lernt

→ **Deep Learning** basiert auf geschichteten neuronalen Netzwerken

Supervised Learning Lernt anhand labeled Daten → predicts unseen data, *Dimensionality Reduction*

Unsupervised Learning unlabeled Daten → Daten-Struktur lernen und Kunden in Zielgruppen teilen (*Recommender Systems*)

Semi-Supervised Learning Mix Sup.+Unsup. → wenig *Labeled* & viele *Unlabeled*

Reinforcement Learning Belohnungsfunktion → sucht ideales Verhalten für maximale Belohnung

Regression Problem Label (vorherzusagendes Attribut) ist kontinuierlich

Classification Problem Label ist kategorial (binary → Bsp. Ja, nein)

Hyperparameters vom Programmierer ausgewählte Parameter (Bsp. K bei K-NN)

Data Quality Assessment

- Identify data sources and trustworthiness
- Check data ranges (e.g. negative salary)
- Interpret statistical key figures
- Validate plausibility of variable correlations
- Measure data redundancy
- Check for anomalies in syntax and semantics
- Explore Null values and duplicate data records

Data Cleaning

Data quality can be improved (although document changes, inform,...):

- identify & remove duplicate data records
- replace null values (e.g. median or mean) or delete them (requires large data)
- investigate the origins of quality issues


Feat. Eng. - Vector Space Model

Data set consists **only** of numerical data (e.g. color → black, red, ...). Categorical data is transformed into numerical data

Statistics

mean μ easy calculate ; median robust against outliers ; $\text{Var}_{\text{simpl.}} \frac{1}{n-1}$; $\text{Var}_{\text{pop.}} \frac{1}{n}$

$\text{Var} \sigma^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu_X)^2$; $\text{Cov}(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu_X)(y_i - \mu_Y)$



Pearson Correlation $\rho(X, Y) = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$ normalizes ranges: -1=anti-corr., +1=corr

Similary Measures

Euclidean Distance $\sqrt{\sum_{i=1}^n (x_i - y_i)^2}$ (distance)

Cosine Distance $1 - \text{CosineSimil.} = 1 - \frac{\sum_{i=1}^n x_i \cdot y_i}{\sqrt{\sum_{i=1}^n x_i^2} \sqrt{\sum_{i=1}^n y_i^2}}$ (distance + angle)

Manhattan Distance $\sum_i |x_i - y_i|$ (along a grid)


Levenshtein/Edit Distance for strings: +1 on delete, add ; +2 on change char

Jacard Similarity: $\frac{|X \cap Y|}{|X \cup Y|} \rightarrow [0, 1]$; 1=same

k-Nearest-Neighbor Classification

Data takes value of majority of k (hyperparameter) neighbors + (optional) weight based on $1/d$ (distance d).

→ very slow with lots of data and most comput. is done in classification stage + least data hungry algo, good baseline for other classifier



Normalization

Min-Max: $\frac{x - \min x}{\max x - \min x} \rightarrow [0, 1]$ + interpretation in %, no negatives - no supervised

Z-Score: $\frac{x - \mu_X}{\sigma_X} \rightarrow \frac{\mu_{X'}}{\sigma_{X'}} = 0, 1$ + for (un)supervised - no direct interpr., negatives

Regression Hypothesis & Evaluation Metrics

Goal is to fit a line or curve-function to a number of points. y_i ground truth, f_i prediction ; **MAE** (Mean Absolute Error), **MAPE** (Mean Absolute Percentage Error) human readable, **MSE** (Mean Squared Error) optimized for processing

$$\frac{1}{m} \left[\sum_{i=1}^m \text{MAE} |y_i - f_i| \quad \sum_{i=1}^m \text{MAPE} \frac{|y_i - f_i|}{y_i} \quad \sum_{i=1}^m \text{MSE} (y_i - f_i)^2 \right] \left| 1 - \frac{\sum_{i=1}^m (y_i - f_i)^2}{\sum_{i=1}^m (y_i - \mu_Y)^2} \right|$$

$R^2 = \rho^2$
 $R^2 \leq 1$; $R^2 < 0$ worse fit than horizontal hyperplane (mean) ; $R^2 > 1$ model performs worse than mean the fraction becomes

Model Evaluation Workflow

Workflow: ① Mix data ② Split into training & test data ③ train ④ evaluate! NEVER reuse test set!

Model Selection: ① 60% Training, 20% Validation, 20% Test ② Train different models ③ Check via Validation and select model ④ Use selection on test data!

K-Fold Cross Valid.: not enough data → ① Split 80% data into k parts ② Round 1 uses first part as **validation** ③ average all accuracy to get final accuracy

Fitting Data

- Underfitting (too loose), Just Right (not perfect fit, but distribution is learnt), Overfitting (too perfect, no distribution)


Regression: X regressor, ϵ error $y_{\text{real}} - y_{\text{lin}} : Y = X\theta + \epsilon$ ($y_j = \theta_0 + \theta_1 \cdot x_j + \epsilon_j$)

Cost Function: used to minimize MSE ; $J(\theta_0, \theta_1) = \frac{1}{N} \sum_{j=N} \epsilon_j^2$

Ordinary Least Squares (OLS): $\frac{\partial J(\theta)}{\partial \theta} = 0 \rightarrow \theta_{\text{opt}} = (X^T X)^{-1} X^T y$ (Gradient Descent iteratively: ① random θ ② small steps down cost surface • direction = negative gradient ; $\theta_{\text{new}} = \theta_{\text{old}} - \alpha \nabla J(\theta) / \partial \theta$)

Regression Performance: $R^2 = 1 - \frac{\sum_j \epsilon_j^2}{\sum_j \epsilon_{\text{OLS}}^2} = 1 - \frac{\sum_j (y_j - \hat{y})^2}{\sum_j (y_j - \bar{y})^2}$

Plot data, as R^2 does not show the distribution (*Anscombe's quartet*)!



To penalize large parameters and stop overfitting, **LASSO** $J(\theta) = \frac{1}{N} \sum_{j=1}^N (y_j - \hat{y}_j)^2 + \lambda \sum_{k=1}^M \theta_k^2$ / $\lambda \sum_{k=1}^M |\theta_k|$ and **Ridge** can be used.

Polynomial Regression: $h(\theta, X) = \theta_0 + \theta_1 X + \theta_2 X^2 + \dots \Rightarrow \theta_0 + \theta_1 X_1 + \theta_2 X_2 + \dots$


Classification

Classification predicts categories (*decision boundary*) vs. **Regression** predicts values (*data relationship*)

Logistic Regression


$$\hat{Y} = P(Y = 1|X) = h_{\theta}(X) = g(X\theta) = \frac{1}{1 + e^{-X\theta}}$$

Instead of either exactly 1 or 0, a percentage $[0, 1]$ is given. No OLS-form possible, good parameters calculated via *Gradient Descent*.



Cross Entropy Cost Function

Classifying for one class **must** penalize the other!



One Versus The Rest

Train model based on a specific class, against others. It is either in that class (1) or not (0) → **One vs. All** classification.

Performance Analysis

$$\text{Sensitivity} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad \text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$
$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}} \quad \text{F1 Score} = \frac{2 \cdot \text{prec} \cdot \text{reca}}{\text{prec} + \text{reca}}$$

	Conf. Matrix		Predict	
	No	Yes	No	Yes
Actual	No	Yes	TN	FP
	Yes	No	FN	TP


Accuracy = $(\text{TP} + \text{TN}) / \text{Total}$; lower F1 indicates some kind of skewed data (unbalance).

Dealing with Bias & Variance

Dealing with bias is easier (underfitting): • larger set of feature • different features • feature engineering • more complex algorithms (more data is not necessarily the answer)

Dealing with variance (overfitting): • reduce features • get more training data • stop training early to avoid overfitting • data cleaning

Bias means we are systematically off by a fixed amount. **Variance** means we are systematically off in a scattered way




Receiver Operating Characteristic (ROC):

- to compare different models
- largest area under curve is best

Decision Trees

Gini Impurity measures the likelihood of an incorrect classification of a random data record. *Impure* features produce no significant information. **Cont. Var.** ① Sort Table ② Calculate avg. of adjacent values (10 → 15 ← 20) ③ Averages are splitting criteria (histogram binning reduces splits)



Label	p(No)	p(Yes)	1-p(No) ² -p(Yes) ²
0%	1	0	0
20%	2/3	1/3	4/9
50%	0	1	0

Gini Impurity of entire feature:

$$2/12 \times 0 + 6/12 \times 4/9 + 4/12 \times 0 = 2/9$$

Regression Trees • Are grown like decision trees for classification but with different splitting criteria • Instead of Gini impurity, **minimize variance of values in the same subset** • Finally, predict average value of all instances in a leaf node

Tree Construction Rules

- Only** positive or negative nodes left → stop branch & assign corresponding decision as leaf node
- Some** positive and negative remain → apply another feature
- No** instances left → look at parent node and decide according the more frequent label
- Instances left but **no** features → training data is no good or contains hidden features ; decide on the more frequent label

Pros & Cons Decision & Regression Trees


+ readable, can both handle numeric & categorical, (almost) no data preparation, perform well on larger datasets - not as accurate, danger of over-complex trees, very sensitive to data quality

Random Forests (collection of decision trees): Utilize **random** (row)sampling with replacement and **random** (column) feature selection (without replacement) to be less sensitive & no overfitting.

Ensembles complex ML model built from simpler models. *Building Bricks* → **Bragging** computes weighted average of votes from simple models, **Boosting** improves simple model by correcting its errors ($F_m(x) + h_m(x)$, e.g. Gradient Boosting).

Neural Networks

Input x_i from previous 'axons', contain some kind of information / **Weights** b, w_i define the inputs' influence / **Sum** \sum sum of the weighted inputs / **Activation** $g(z)$ can be **hard threshold** ($\alpha = g(z) = 1$ if $\sum \geq 0$ else 0), **logistic regression** ($g(z) = \frac{1}{1+\exp(-z)}$)



Feed Forward data moves one way (left → right)

Hidden Layers: allows for new features a to be added

Layer Count $c_L = c_{\text{hidden}} + c_{\text{output}}$ (input is not counted)

Feed Forward →	Back Propagation ←	Gradient Descent ↓
<ul style="list-style-type: none">A neural net is the connection of many artificial neurons in parallel and seriesEach layer of a feed-forward ANN is defined by its weights W and biases bGiven data with labels y and features x, we feed x forward through the ANN	<ul style="list-style-type: none">At the output of the ANN we get the estimate \hat{y} of yBecause we start with random W and b, the estimate \hat{y} is poorDefine the error J as a measure of distance between \hat{y} and yWe seek the best W & b to minimize J. This occurs when the gradient of J is small or zeroSo we need the gradient of JCompute the gradient iteratively at each layer: more efficient!We introduce deltas δ at each layer to simplify computation	<ul style="list-style-type: none">Back-propagation gives the gradient of J with respect to weights W and biases bThe gradient is a vector which points steeply up the error surfaceSo take a small step α in the opposite direction, reducing J"Take a step" means making small changes to W & bChanges are defined exactly by $-\alpha$ times the gradientThis is gradient descent by feed forward & back-prop using labelled data

Activation Function

ReLU (rectified linear unit ;good for hidden) $\max(0, x)$ brings better performance and alleviates the vanishing gradients problem (flat line activation) → introduces **dying units** (bad for gradient training)




Leaky ReLU $\max(0.1x, x)$ introduces non-zero for negative values and fixes vanishing gradients + dying unit → results aren't always consistent

Soft-Max (good for output) $\text{softmax}_j = \frac{\exp(z_j)}{\sum_{k=1}^K \exp(z_k)}$ is a generalization of the logistic function to **multiple dimensions** and can be interpreted as probabilities (normalization).

Convolution Neural Network

Features are extracted from images using **convolution** (Faltung). **Weights** are used as Masks/Filters on these images to extract certain features such as edges or colour. The final layer mostly uses **softmax**!



Pooling applies operations such as max, min or avg to a pixel neighbour and the neighbourhood size is determined by the pool size parameter.

Parameter Counting Example on 20×20 -RGB-Image (1) \geq *Filters of 3×3 ($(3 \cdot 3 \cdot 3 + 1) \cdot 2 = 56$ results $18 \times 18 \times 2$ Layers)* (2) \geq *Filters of 3×3 with 2 channels each ($(3 \cdot 3 \cdot 2 + 1) \cdot 3 = 57$ results $16 \times 16 \times 3$ Layers)* (3) *Flatten result to a single column of neurons ($1 \cdot 16 \cdot 16 \cdot 3 = 1 \times 768$)* (4) *Fully connected (aka. dense) layer and 2 output channels (#param = $768 \cdot 2 + 2 = 1538$ (2 biases! & 2 neurons/output shape))*

Natural Language Processing

Syntax compare spelling or words / **Similarity** Levenshtein Distance

Word Relatedness Words are *related*... • when they appear in the **same document**

(synonyms not related!!) • when they relate to the **same topic** (requires a lot of space)

• when they can be replaced by each other

Word Similarity Words are *similar* when they appear in the **same context**

Term Frequency (Inverse Document Frequency) $TF\text{-}IDF(x, y) = TF(x, y) \cdot \log\left(\frac{N}{DF(x)}\right)$

/ N number of documents in the corpus, TF counts the number of occurrences of x in y , DF counts the number of documents that contain x . *QUETZALCOATLUS* has high score, due to being used a lot in a **single** document, compared to the low scoring word *AND* (used everywhere).

Recurrent Neural Nets (RNN)


$h_t = (W_{xh} \cdot x_t + W_{hh} \cdot h_{t-1} + b_h) \quad y_t = \varphi(W_{hy} \cdot h_t + b_y)$

RNN are made for the supervised analysis of arbitrary time-series data (instead of fixed length, various sized strings can be analyzed) and **do not parallelize**! LSTM & GRU help with alleviating vanishing gradients to remember more of the context (Training an RNN by multiplying gradients backwards in time).

GRU (~~h_t~~) use hidden states as memory (faster) & **Long Short-Term Memory Model** have an additional memory state (better performance)

Bidirectional RNN reads left and right → produces together output

Encoder-Decoder Architecture: usually RNNs have one input and output, which makes it unusable for i.e. translators. To allow in- and outputs of various lengths, en- & decoder architecture is used.



Attention Mechanism

Mimics the behaviour of knowing which page the content can be found in. Words are assigned an *importance* value which is used by the next generation for the output words.



k-Means Clustering

- ① Input k Clusters and data points x_n
- ② Randomly choose k clusters centers μ_k
- ③ Repeat until convergence (stability): (a) assign data points to nearest cluster center (b) update cluster center with mean of assigned datapoints

Two data points can meet in a cluster and separate again in the next step! → pattern always changes

Clustering Distortion

Total distortion sum of squared distances of points and cluster center • **Average Distortion** average of total distortion (additional $1/n$) and allows for comparison with other data sets

k-Means only approximates optimal solutions and converges either to global **OR** local minimum


Elbow Methods increasing the amount of clusters leads to an 'elbow' in the distance-nCluster diagram → use cluster count at the 'elbow'

Agglomerative Clustering

Outputs an hierarchy (*dendrogram*) that are informative and human interpretable. It does not scale with larger data, but cluster stays together forever! ① Initially each datapoint is a separate cluster ② While stop condit. (on no change) is not met: (a) calculate distance between pairs of cluster (b) merge pairs with shortest distance

Association Rules

Implication ($X \rightarrow Y$): the higher confidence, the more likely Y is also present in transaction with X



Support (Interestingness): how frequently an item set occurs in the data

$\text{support}(\{i_1, \dots, i_n\}) = \frac{\# \text{purchases of } \{i_1, \dots, i_n\}}{\# \text{transactions}}$

Support of an Association $\text{support}(X \rightarrow Y) = \text{support}(Y \rightarrow X) = \text{support}(X \cup Y)$

Confidence (Trustworthiness): how frequently items in Y appear in transactions that contain X

$\text{confidence}(X \rightarrow Y) = \frac{\text{support}(X \cup Y)}{\text{support}(X)}$

BOTH Support and Confidence need to be considered. — Only Antecedent X and co-occurrence $X \cup Y$ is taken into account, not Y ; sup. & conf. cannot filter coincidental occurrences out.

Lift (Association Strength): how many times more often X and Y show up together than expected

$\text{lift}(X \rightarrow Y) = \frac{\text{support}(X \cup Y)}{\text{support}(X) \cdot \text{support}(Y)}$ (= 1 independent; < 1 X & Y appear less often together ; > 1 X & Y appear often together)

Apriori Algorithm

- ① Generate frequent item sets → ① frequent items leads subsets being frequent too ② infrequent leads supersets infrequent too. (shape **lattice of subsets**: first layer single node, second double connection,...)

- ② Extract rules from frequent item satisfying the confidence threshold → ① when an item violates the confidence bound all subsets can be eliminated. (shape **lattice of rules**)

Recommender System - Non-Personalized

+ simple, fast & cheap ; Associations are context aware ; Associations can be tailored to specific businesses (who viewed, who bought...) ; No new user cold start problem — non-personalized

Recommender System - Personalized

$$P_{uj} = \frac{\sum_{i \in N_j} \text{similar}(j, i) \cdot r_{ui}}{\sum_{i \in N_j} \text{similar}(j, i)} \quad P_{uj} = \bar{r}_u + \frac{\sum_{i \in N_u} \text{similar}(u, i) \cdot (r_{ij} - \bar{r}_i)}{\sum_{i \in N_u} \text{similar}(u, i)}$$

N rated items (hyperparameter) ; P_{uj} how much does user u like item j ; N_j Neighbours of item j ; r_{ui} rating that user u gave item i ; N_u Neighbors of user u ; r_{ij} from neighbor i to item j

Personalized Content-Based Recommendations (left equ.): + user independence, no new item cold start problem — limited content analysis, overspecialization (bubble esque), severe new user cold start problem

User-to-User Collaborative Filtering (right equ.): + cross-category recommendations — many participants required, system must first learn user + item rating (cold start problem), gray sheep problem (unusual taste results in low accuracy)

Hybridization

Combination of different recommender systems! • **Weighted Approach** item score is weighted sum score from different recommenders • **Mixed Approach** results of different recommender are presented • **Cascade Approach** one recommender refines the result of other method • **Switching Approach** use different methods ta different places (online shops)